

practice2_raw

October 27, 2023

1 Práctica 2: PCA - Red Wine

Román Alberto Vélez Jiménez

CU: 165462

De la base de datos [red wine](#) de kaggle, se hará un análisis PCA.

```
[ ]: # imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from mpl_toolkits.mplot3d import Axes3D
import plotly.graph_objs as go
import plotly.express as px
```

2 Parte 1: Data

2.1 Lectura de Datos

```
[ ]: # read data from numpy
red_wine = pd.read_csv("data/winequality-red.csv")
red_wine
```

```
[ ]:      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
0              7.4             0.700         0.00         1.9         0.076
1              7.8             0.880         0.00         2.6         0.098
2              7.8             0.760         0.04         2.3         0.092
3             11.2             0.280         0.56         1.9         0.075
4              7.4             0.700         0.00         1.9         0.076
...           ...             ...             ...             ...             ...
1594            6.2             0.600         0.08         2.0         0.090
1595            5.9             0.550         0.10         2.2         0.062
1596            6.3             0.510         0.13         2.3         0.076
1597            5.9             0.645         0.12         2.0         0.075
1598            6.0             0.310         0.47         3.6         0.067
```

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	11.0	34.0	0.99780	3.51	0.56
1	25.0	67.0	0.99680	3.20	0.68
2	15.0	54.0	0.99700	3.26	0.65
3	17.0	60.0	0.99800	3.16	0.58
4	11.0	34.0	0.99780	3.51	0.56
...
1594	32.0	44.0	0.99490	3.45	0.58
1595	39.0	51.0	0.99512	3.52	0.76
1596	29.0	40.0	0.99574	3.42	0.75
1597	32.0	44.0	0.99547	3.57	0.71
1598	18.0	42.0	0.99549	3.39	0.66

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
...
1594	10.5	5
1595	11.2	6
1596	11.0	6
1597	10.2	5
1598	11.0	6

[1599 rows x 12 columns]

```
[ ]: # drop 'quality'
quality_by_wine = red_wine['quality'].copy()
quality_by_wine = 'quality_' + quality_by_wine.astype(str)
red_wine = red_wine.drop('quality', axis='columns')
```

2.2 Preparando los Datos

```
[ ]: # get numpy array
X = red_wine.copy().to_numpy()
X.shape
```

```
[ ]: (1599, 11)
```

```
[ ]: # scale array
Xstand = X - np.mean(X, axis=0)
Xstand = np.divide(Xstand, np.std(Xstand, ddof=1, axis=0))
```

```
[ ]: # look if the data is centered
print(f"data is centered: {np.all(np.isclose(Xstand.mean(axis=0), 0))}")
```

```
# look if the data is centered
print(f"data is scaled: {np.all(np.isclose(Xstand.std(axis=0, ddof=1), 1))}")
```

data is centered: True


data is scaled: True

```
[ ]: # variance-covariance matrix
m_size = Xstand.shape[0]
S = (1 / (m_size - 1)) * (Xstand.T @ Xstand)
```

Intel MKL WARNING: Support of Intel(R) Streaming SIMD Extensions 4.2 (Intel(R) SSE4.2) enabled only processors has been deprecated. Intel oneAPI Math Kernel Library 2025.0 will require Intel(R) Advanced Vector Extensions (Intel(R) AVX) instructions.

2.3 Mapa de Calor

```
[ ]: # create a heatmap
plt.imshow(S, cmap='PiYG', interpolation='nearest')
plt.colorbar()
# add ticks
plt.xticks(range(len(red_wine.columns)), red_wine.columns, rotation=90)
plt.yticks(range(len(red_wine.columns)), red_wine.columns, rotation=0)
# add title
plt.title("Mapa de calor de la matriz $$$")
# Display the plot or save it to a file
plt.show()
```



3 Parte 2: PCA

3.1 Cálculo de eigenvectores de $S = A^T A$

```
[ ]: # get eigenvectors
eigvalues, eigvectors = np.linalg.eig(S)
```

3.2 Proyección Datos

```
[ ]: # proyect data
Xstand_proy = Xstand @ eigvectors
```

4 Parte 3: Varianza Acumulada

```
[ ]: import numpy as np
import matplotlib.pyplot as plt

# Assuming eigvalues is a numpy array containing the eigenvalues
# Sort the eigenvalues in descending order
eigvalues.sort()
eigvalues = eigvalues[::-1]

# Calculate the cumulative explained variance
cumulative_variance = eigvalues.cumsum() / eigvalues.sum()

# Create a figure and axis for the cumulative variance explained graph
fig, ax = plt.subplots()

# Plot the cumulative explained variance
ax.plot(np.arange(1, eigvalues.size + 1), cumulative_variance, '-o')


# Set better formatting for the tick labels
ax.set_xticks(np.arange(1, eigvalues.size + 1))
ax.set_xticklabels(np.arange(1, eigvalues.size + 1))
ax.set_xlabel("Numero de Componentes")
ax.set_ylabel("Varianza Explicada")
ax.grid(True)

# Annotate data points with y-values
for i, var in enumerate(cumulative_variance):
    ax.text(i + 1, var, f'{var:.2f}', ha='right', va='bottom')

# Add a point at (0,0) and label it on the cumulative variance explained graph
ax.plot(0, 0, 'ro') # Red dot at (0,0)
ax.text(0, 0, "(0, 0)", ha='right', va='bottom')

# Create a bar chart for individual contributions
individual_contributions = eigvalues / eigvalues.sum()
x = np.arange(1, eigvalues.size + 1)
ax2 = ax.twinx() # Create a secondary y-axis
ax2.bar(x, individual_contributions, align='center', alpha=0.5, color='g',
        label='Contribución Individual')
ax2.set_ylabel("Contribution to Variance")

# Show the plot
plt.title("Contribución de cada Componente")
plt.tight_layout()
plt.show()
```



5 Parte 4: Contribución

Doubt: does the PC is the column of eigenvecs or the rows?

```
[ ]: # look eigenvectors
df_eig = pd.DataFrame(
    eigvectors,
    columns=['PC' + str(i+1) for i in range(eigvalues.size)],
    index=red_wine.columns
)
df_eig
```

```
[ ]:
```

	PC1	PC2	PC3	PC4	PC5	\
fixed acidity	0.489314	-0.110503	-0.123302	-0.229617	-0.082614	
volatile acidity	-0.238584	0.274930	-0.449963	0.078960	0.218735	
citric acid	0.463632	-0.151791	0.238247	-0.079418	-0.058573	
residual sugar	0.146107	0.272080	0.101283	-0.372793	0.732144	
chlorides	0.212247	0.148052	-0.092614	0.666195	0.246501	
free sulfur dioxide	-0.036158	0.513567	0.428793	-0.043538	-0.159152	
total sulfur dioxide	0.023575	0.569487	0.322415	-0.034577	-0.222465	

density	0.395353	0.233575	-0.338871	-0.174500	0.157077
pH	-0.438520	0.006711	0.057697	-0.003788	0.267530
sulphates	0.242921	-0.037554	0.279786	0.550872	0.225962
alcohol	-0.113232	-0.386181	0.471673	-0.122181	0.350681

	PC6	PC7	PC8	PC9	PC10	\
fixed acidity	-0.639691	-0.249523	0.194021	-0.177595	-0.350227	
volatile acidity	-0.002389	0.365925	-0.129110	-0.078775	-0.533735	
citric acid	0.070910	0.621677	-0.381450	-0.377516	0.105497	
residual sugar	-0.184030	0.092872	0.007523	0.299845	0.290663	
chlorides	-0.053065	-0.217671	0.111339	-0.357009	0.370413	
free sulfur dioxide	0.051421	0.248483	0.635405	-0.204781	-0.116596	
total sulfur dioxide	-0.068702	-0.370750	-0.592116	0.019036	-0.093662	
density	0.567332	-0.239990	0.020719	-0.239223	-0.170481	
pH	-0.340711	-0.010970	-0.167746	-0.561391	-0.025138	
sulphates	-0.069555	0.112320	-0.058367	0.374604	-0.447469	
alcohol	0.314526	-0.303015	0.037603	-0.217626	-0.327651	

	PC11
fixed acidity	0.101479
volatile acidity	0.411449
citric acid	0.069593
residual sugar	0.049156
chlorides	0.304339
free sulfur dioxide	-0.014000
total sulfur dioxide	0.136308
density	-0.391152
pH	-0.522116
sulphates	-0.381263
alcohol	0.361645

```
[ ]: # for each component, get the index with highest value
list_max_contr = []
for name, pc in df_eig.T.iterrows():
    # get maximum contributor
    pc = pc.sort_values(key=np.abs, ascending=False) # sort by abs value
    max_value = pc[0]
    max_contrib_var = pc.index[0]
    # save
    list_max_contr.append((name, max_contrib_var, max_value))

# to dataframe
pd.DataFrame(list_max_contr, columns=['PC', 'variable', 'contribution']).
    ↪set_index('PC')
```

```
[ ]:          variable  contribution
PC
```

PC1	fixed acidity	0.489314
PC2	total sulfur dioxide	0.569487
PC3	alcohol	0.471673
PC4	chlorides	0.666195
PC5	residual sugar	0.732144
PC6	fixed acidity	-0.639691
PC7	citric acid	0.621677
PC8	free sulfur dioxide	0.635405
PC9	pH	-0.561391
PC10	volatile acidity	-0.533735
PC11	pH	-0.522116

```
[ ]: # look S matrix
pd.DataFrame(
    S,
    columns=red_wine.columns,
    index=red_wine.columns
)
```

```
[ ]:
fixed acidity    fixed acidity    volatile acidity    citric acid \
fixed acidity    1.000000        -0.256131        0.671703
volatile acidity  -0.256131        1.000000       -0.552496
citric acid       0.671703        -0.552496        1.000000
residual sugar    0.114777        0.001918        0.143577
chlorides         0.093705        0.061298        0.203823
free sulfur dioxide -0.153794       -0.010504       -0.060978
total sulfur dioxide -0.113181       0.076470        0.035533
density           0.668047        0.022026        0.364947
pH               -0.682978        0.234937       -0.541904
sulphates         0.183006       -0.260987        0.312770
alcohol          -0.061668       -0.202288        0.109903
```

```

fixed acidity    residual sugar    chlorides    free sulfur dioxide \
fixed acidity    0.114777    0.093705        -0.153794
volatile acidity  0.001918    0.061298       -0.010504
citric acid       0.143577    0.203823       -0.060978
residual sugar    1.000000    0.055610        0.187049
chlorides         0.055610    1.000000        0.005562
free sulfur dioxide 0.187049    0.005562        1.000000
total sulfur dioxide 0.203028    0.047400        0.667666
density           0.355283    0.200632       -0.021946
pH               -0.085652   -0.265026        0.070377
sulphates         0.005527    0.371260        0.051658
alcohol           0.042075   -0.221141       -0.069408
```

```

fixed acidity    total sulfur dioxide    density    pH    sulphates \
fixed acidity    -0.113181    0.668047   -0.682978    0.183006
```


volatile acidity	0.076470	0.022026	0.234937	-0.260987
citric acid	0.035533	0.364947	-0.541904	0.312770
residual sugar	0.203028	0.355283	-0.085652	0.005527
chlorides	0.047400	0.200632	-0.265026	0.371260
free sulfur dioxide	0.667666	-0.021946	0.070377	0.051658
total sulfur dioxide	1.000000	0.071269	-0.066495	0.042947
density	0.071269	1.000000	-0.341699	0.148506
pH	-0.066495	-0.341699	1.000000	-0.196648
sulphates	0.042947	0.148506	-0.196648	1.000000
alcohol	-0.205654	-0.496180	0.205633	0.093595

	alcohol
fixed acidity	-0.061668
volatile acidity	-0.202288
citric acid	0.109903
residual sugar	0.042075
chlorides	-0.221141
free sulfur dioxide	-0.069408
total sulfur dioxide	-0.205654
density	-0.496180
pH	0.205633
sulphates	0.093595
alcohol	1.000000

6 Parte 5: Biplot

6.1 3D

```
[ ]: # Create a DataFrame with your data
from turtle import title

df = pd.DataFrame(Xstand_proy[:, :3], columns=["PC1", "PC2", "PC3"])
df["quality"] = pd.Categorical(quality_by_wine, categories=quality_by_wine.
    ↪sort_values().unique(), ordered=True)
df = df.sort_values('quality', ascending=False)

# Create an interactive 3D scatter plot with the "Viridis" color scale
fig = px.scatter_3d(df, x="PC1", y="PC2", z="PC3", color="quality",
    ↪title="Primeras 3 Componentes Principales", color_discrete_sequence=px.
    ↪colors.sequential.Viridis_r)

# Add a legend
fig.update_traces(marker=dict(size=3), selector=dict(mode='markers'),
    ↪showlegend=True)

# Show the plot
```

```
fig.show()
```

6.2 2D

```
[ ]: df['quality'].cat.codes.unique()
```


```
[ ]: array([2, 3, 4, 1, 5, 0], dtype=int8)
```

```
[ ]: # Scatter plot of the first two vectors
scatter = plt.scatter(df['PC1'], df['PC2'], c=df['quality'].cat.codes,
    ↪ cmap='viridis', s=5)

# Add labels
plt.xlabel('PC1')
plt.ylabel('PC2')

# Create a legend
quality_values = df['quality'].sort_values().cat.codes.unique()
legend_elements = [plt.Line2D([0], [0], marker='o', color='w',
    ↪ markerfacecolor=scatter.cmap(scatter.norm(quality_values[i])),
                                markersize=10, label=quality_values[i]+3) for i
    ↪ in range(len(quality_values))]
plt.legend(handles=legend_elements, title='Quality')

# Show the plot
plt.suptitle("Datos Proyectados")
plt.title("Primeras 2 Componentes Principales")
plt.tight_layout()
plt.show()
```



[]:

[]:

[]:

[]: